

# Course Introduction: Operating Systems

Mainack Mondal and Saptarshi Ghosh

CS31202 / CS30002

Spring 2022-23



# Today's class

- Resources and notes
- Introduction to operating systems

# Instructors



- **Mainack Mondal**
  - Research interests: system security, large scale system measurement, privacy
  - Office: CSE 316
  - Will also be in your lab

# Instructors



- Saptarshi Ghosh
  - Research interests: Social Networks, ML, NLP, IR
  - Office: CSE 207
  - Will also be in your lab

# Teaching assistants

- Soham Poddar (PhD)
- Shounak Paul (PhD)
- Shubham Soni (M.Tech.)
- K Ganesh Reddy (M.Tech.)
- D Vamshidhar Reddy (Dual)
- Rahul Aditya (Dual)
- Karnam Sai Keerthana (Dual)
- Aayush Prasad (Dual)
- B Smayan Das (Dual)
- Aayushi Vidyanta (Dual)

# Requirements

- Using computers and application software

# Requirements

- Using computers and application software
- How to write, compile and run C-programs

# Requirements

- Using computers and application software
- How to write, compile and run C-programs
- Prerequisites
  - Computer Organization and architecture
  - Algorithms and Data Structure



# Website / Books

- Website:

<https://cse.iitkgp.ac.in/~mainack/courses/2022-spring/OS-course/index.html>

- Textbooks / References:

- Operating Systems Concepts, 9<sup>th</sup> ed. - A. Silberschatz, P.V. Galvin, and G. Gagne. Wiley
- Some recent topics from research papers

# Course logistics

- End-semester exam – 50%
- Mid-semester exam – 30%
- One class test – 20%

# Why should you attend class?

- Slides are not complete and exam questions will come from class coverage
- Tutorial and doubt clearing sessions

# Last but not the least

- Ask questions in the class
  - You need to know how Operating Systems (OS) as computer scientist
  - It is best done via class - interaction

# Last but not the least

- Ask questions in the class
  - You need to know how Operating Systems (OS) as computer scientist
  - It is best done via class - interaction
- Do lab assignments & practice problems religiously
  - This is a systems course
  - There is no other way of learning how systems work than actually trying to build systems

# Today's class

- Resources and notes
- Introduction to operating systems

# Today's class

- Resources and notes
- Introduction to operating systems (OS)
  - What is an OS
  - What are the goals of an OS
  - Under the hood: the structure of OS
  - How does OS work?

**What is an Operating System?**



Let's start with enumerating the actors

# Let's start with enumerating the actors

Users

might be one or multiple (using servers)

# Let's start with enumerating the actors

Users

might be one or multiple (using servers)

Logs in and start applications

gcc, gedit, notepad, firefox, chrome, excel, ...

# Let's start with enumerating the actors

Users

might be one or multiple (using servers)

Logs in and start applications

gcc, gedit, notepad, firefox, chrome, excel, ...

These applications ultimately need to use hardware

processor, ram, display, mouse, keyboard ....

# Let's start with enumerating the actors

Users

might be one or multiple (using servers)

Logs in and start applications

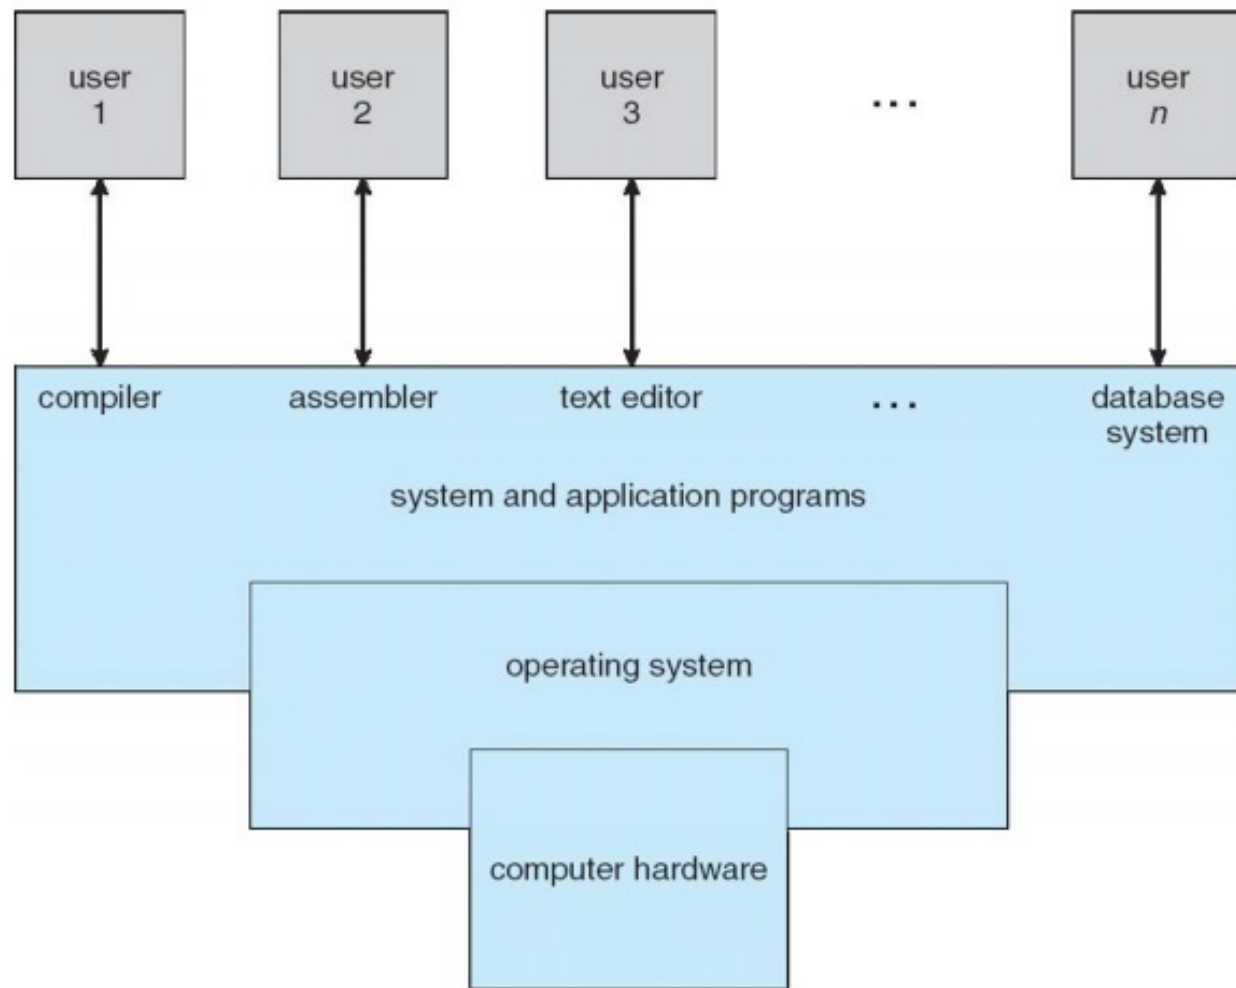
gcc, gedit, notepad, firefox, chrome, excel, ...

These applications ultimately need to use hardware

processor, ram, display, mouse, keyboard ....

**Operating system seats right in between applications  
and hardware**

# Putting it all together



# OS: Definition

A program that acts as an intermediary between users of a computer system and hardware

# Today's class

- Resources and notes
- Introduction to operating systems (OS)
  - What is an OS
  - What are the goals of an OS
  - Under the hood: the structure of OS
  - How does OS work?



# Goals of an Operating System

# Primary goals

- Make the computer system convenient and easy to use (for the users)
- Ensure efficient utilization of resources (processor time, printer, RAM, ...)
  - Controls and co-ordinates the use of hardware resources among multiple users and applications

These goals might be conflicting  
across computing systems

# These goals might be conflicting across computing systems

Mini computers  
and mainframes  
(1970 - 80)

Best utilization  
of resources

Cost: user  
experience  
(often no  
immediate  
result)

# These goals might be conflicting across computing systems

Mini computers  
and mainframes  
(1970 - 80)

Workstations /  
terminals  
(1980 - now)

Best utilization  
of resources

Fast response  
time

Cost: user  
experience  
(often no  
immediate  
result)

Cost: Resource  
utilization is  
suboptimal

# These goals might be conflicting across computing systems

Mini computers  
and mainframes  
(1970 - 80)

Workstations /  
terminals  
(1980 - now)

Laptops

Best utilization  
of resources

Fast response  
time

Flexible UI, easy  
to use

Cost: user  
experience  
(often no  
immediate  
result)

Cost: Resource  
utilization is  
suboptimal

Cost: single user  
system

# These goals might be conflicting across computing systems

Mini computers  
and mainframes  
(1970 - 80)

Workstations /  
terminals  
(1980 - now)

Laptops

Mobile systems

Best utilization  
of resources

Fast response  
time

Flexible UI, easy  
to use

Optimized usability  
& battery life

Cost: user  
experience  
(often no  
immediate  
result)

Cost: Resource  
utilization is  
suboptimal

Cost: single user  
system

Cost: Severely  
constrained  
resources (your  
android cannot  
handle running full  
fledged game and  
browsing)

# These goals might be conflicting across applications

Application A: I need 3 GB of memory, now!



# These goals might be conflicting across applications

Application A: I need 3 GB of memory, now!

Application B: I need 2 GB of memory, now!

# These goals might be conflicting across applications

Application A: I need 3 GB of memory, now!

Application B: I need 2 GB of memory, now!

Hardware: oops, total RAM is 4 GB

# OS: goal-oriented definition

- A resource allocator
  - Manages all resources (processor time, RAM, display, ...) to ensure they are shared in an **efficient** and **fair** manner
  - Decides which application gets how much resources and when
- A control program
  - Controls execution of other programs / applications to prevent errors and improve usability (e.g., by giving faster response to users)
  - A faulty application should not disrupt other applications (or the OS itself)

# OS: goal-oriented definition (contd.)

- Enabler of communication / coordination
  - One application may need to communicate to others, or share data / state / ...
- Enabler of easier development of applications
  - Offers a set of common services for applications (e.g., for I/O) - application developer does not need to worry about specifics of devices
  - Gives an illusion of infinite resources – dedicated processor, infinite memory, ...

# Today's class


- Resources and notes
- Introduction to operating systems (OS)
  - What is an OS
  - What are the goals of an OS
  - Under the hood: the structure of OS
  - How does OS work?

# Typical structure of OS (unix)

- Note: By structure we simply meant control flow


# Typical structure of OS (unix)

- Note: By structure we simply meant control flow



# Typical structure of OS (unix)


- Note: By structure we simply meant control flow






# Typical structure of OS (unix)

- Note: By structure we simply meant control flow



# Typical structure of OS (unix)

- Note: By structure we simply meant control flow




# Function of Kernel

- Part of OS closest to the hardware, handles important functionalities
  - Managing memory, network, file, processes, system calls ...

# Function of Kernel

- Part of OS closest to the hardware, handles important functionalities
  - Managing memory, network, file, processes, system calls ...
- Other parts of OS (e.g., the shell) and application programs can interact with kernel whenever they require these functionalities
  - E.g., need to read from keyboard (*scanf*), show something on display (*printf*), create directory (*mkdir*)

# What other services?



Source: Silberschatz, Galvin and Gagne ©2013

# Food for thought

- A user might ask for *printf* any time
  - Does “Kernel run all the times”?

# Food for thought

- A user might ask for *printf* any time
  - Does “Kernel run all the times”?
- Possible problem
  - Kernel needs resources to run

# Food for thought

- A user might ask for *printf* any time
  - Does “Kernel run all the times”?
- Possible problem
  - Kernel needs resources to run
- Insight
  - Kernel “only” needs to run when any of its functionalities are required



# Today's class

- Resources and notes
- Introduction to operating systems (OS)
  - What is an OS
  - What are the goals of an OS
  - Under the hood: the structure of OS
  - How does OS work?

# Today's class

- Resources and notes
- Introduction to operating systems (OS)
  - What is an OS
  - What are the goals of an OS
  - Under the hood: the structure of OS
  - How does ~~OS~~ work?  
Kernel

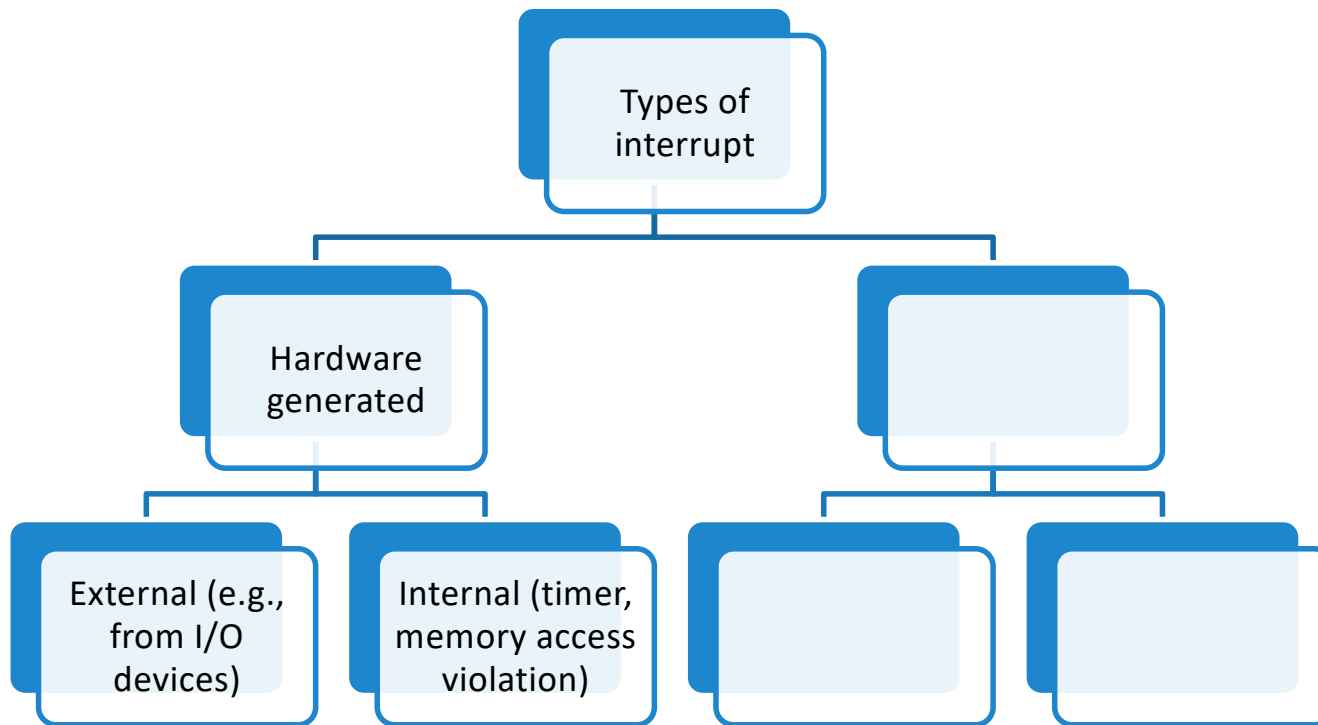
# Kernel is “interrupt” driven

- Think of interrupt as the “wake up call” to kernel
  - When interrupt comes, some function in Kernel is invoked

# Kernel is “interrupt” driven

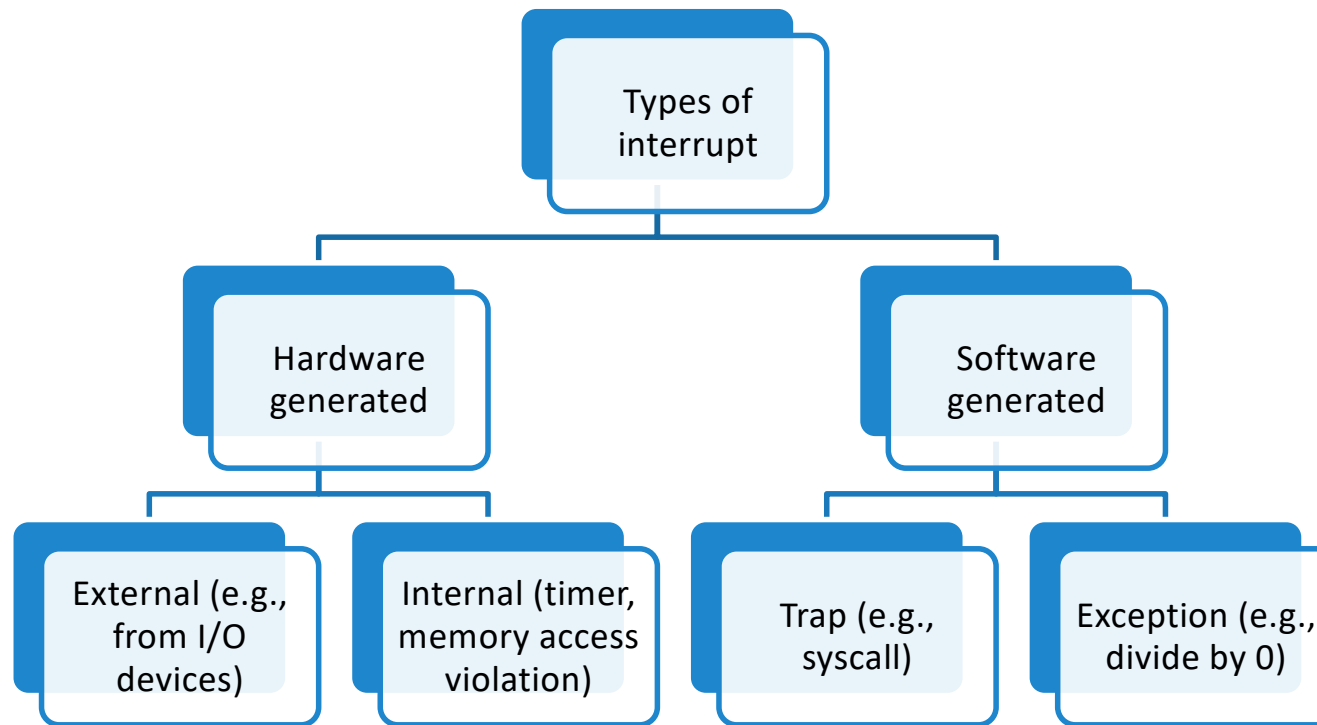
- Think of interrupt as the “wake up call” to kernel
  - When interrupt comes, some function in Kernel is invoked
- More technically
  - Interrupt is a signal (instruction), generated by hardware/software
  - The interrupt in turn activates appropriate kernel routine(s) depending on specific category of the interrupt

# Types of interrupt



Hardware interrupts: generated by external device (e.g., a printer)  
OR internal hardware unit

# Types of interrupt




Hardware interrupts: generated by external device (e.g., a printer)  
OR internal hardware unit

Software interrupts: Generated by a running program

# **Case study: Handling Input/Output (I/O) requests**


# The setup



Source: Silberschatz, Galvin and Gagne ©2013




# The setup (contd.)



- I/O devices and CPU can execute in parallel
  - Each device controller is in charge of a particular device type
  - Each device controller has a local buffer
  - Data transfers between local buffer and main memory
  - Device controller sends an interrupt to the CPU to indicate I/O completion


# The setup (contd.)



**Question: How would an I/O request be handled in this setup?**


- Each device controller has a local buffer
- Data transfers between local buffer and main memory
- Device controller sends an interrupt to the CPU to indicate I/O completion

# Handling I/O in the setup




- Option 1
  - User program makes I/O request - transfers data from memory to buffer of device controller and initiates I/O
  - CPU remains idle as long as device controller handles request
  - Upon I/O completion, device controller sends interrupt, then execution of user program resumes
  - "Busy waiting" - under-utilization of CPU

# Handling I/O in the setup




- Option 2: use OS
  - User program makes I/O request through a "system call"
  - OS transfers data from memory to buffer of device controller and initiates I/O
  - OS can allocate CPU to other programs during I/O of one program
  - Upon I/O completion, device controller sends interrupt, then OS may resume this particular user program (according to some scheduling policy)
  - "Multiprogramming" - much better utilization of CPU

# Handling I/O in the setup



- Previous scheme of interrupt-driven I/O is fine for small amounts of data, but high overhead for bulk data transfer such as disk I/O
- Hence option 3 – Direct Memory Access (DMA)
  - Once OS initiates I/O, DMA manages transfer of data between memory and device controller with no intervention by the CPU

# I/O handling in von Neumann architecture



Source: Silberschatz, Galvin and Gagne ©2013