

# Process (contd.)

Saptarshi Ghosh and Mainack Mondal

CS31202 / CS30002



# So far on processes

- What is a process?
- Structure of a process in memory
- Process control block (PCB)
- Process states
- Context switch

**Process queues (for scheduling)**

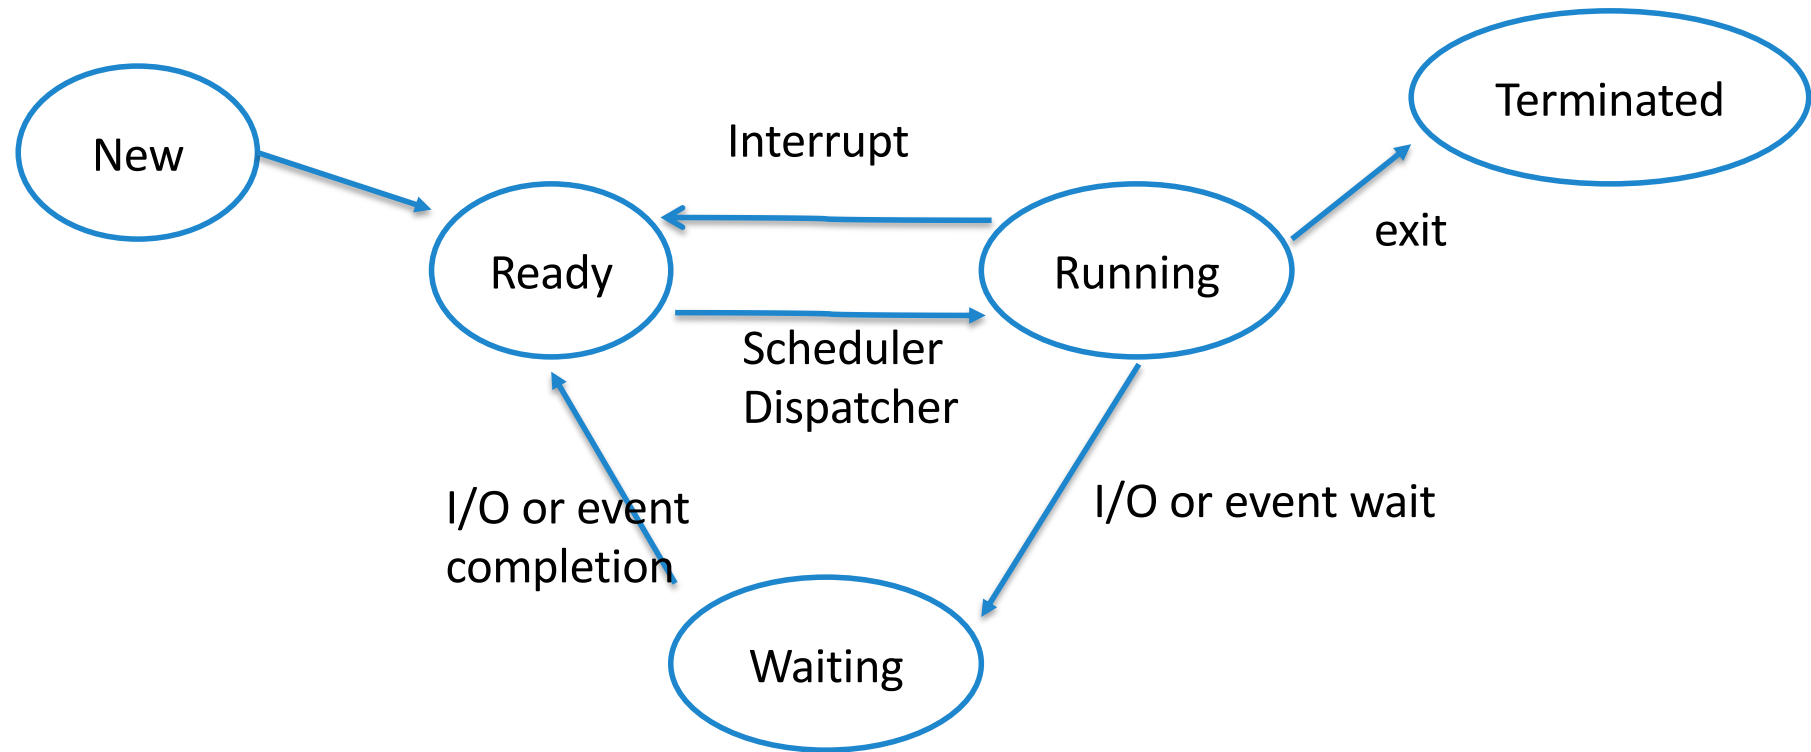
# Process queues for scheduling

- Several scheduling queues exist in OS
  - A PCB is linked to one of the queues at any given time
  - The PCBs in a queue are connected as a linked list

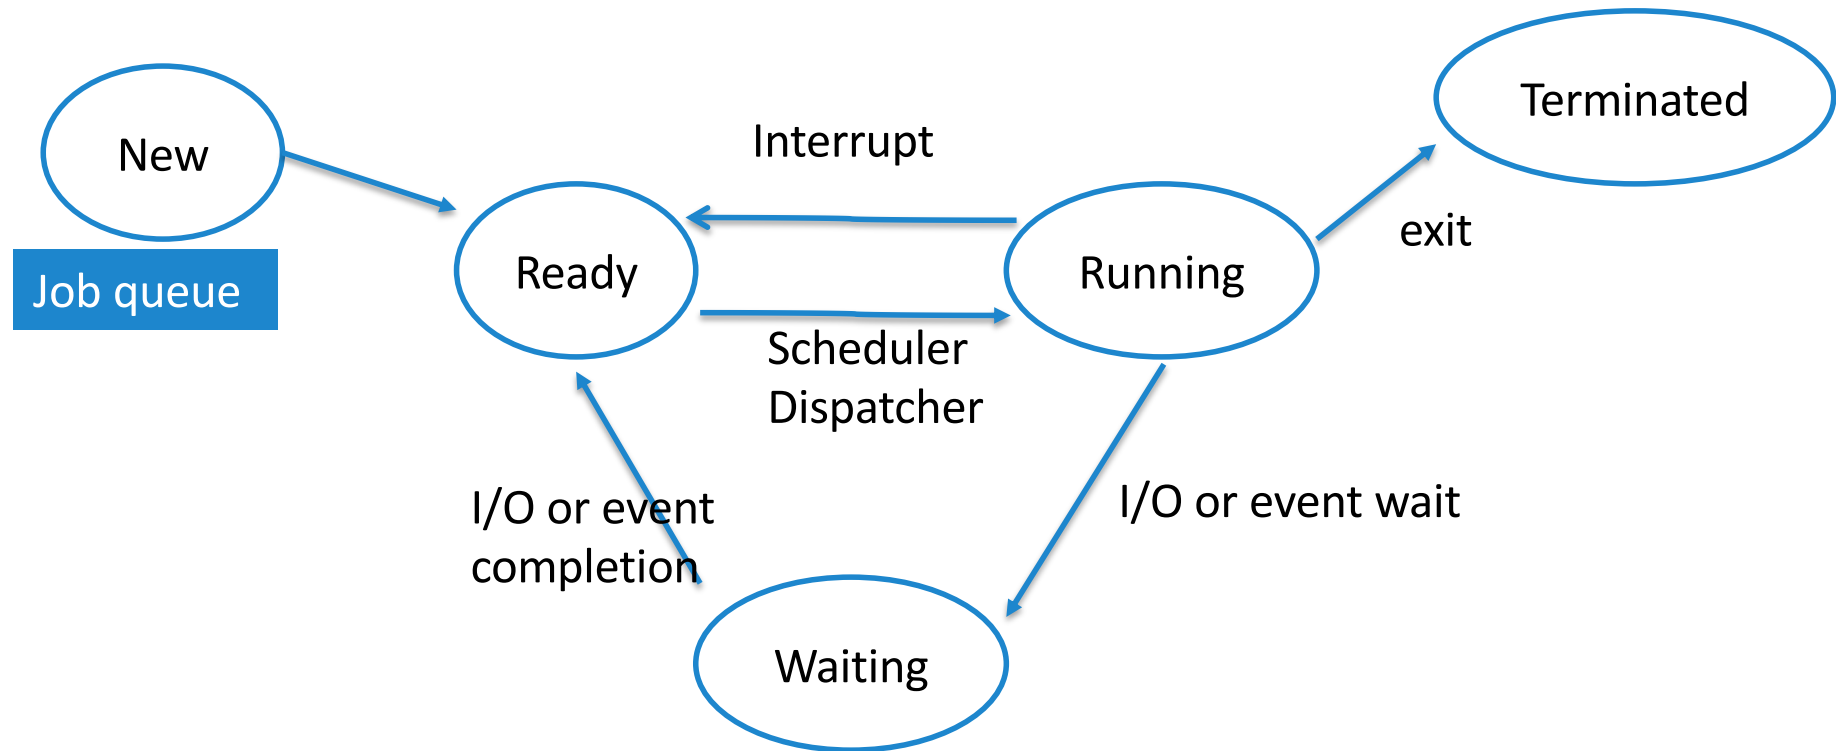
# Characteristics of process queues

- Each I/O device has its own device queue
  - Contains (PCBs of) processes waiting for this device
- Each event also has its own queue
  - Contains processes waiting for this event
- Process scheduling can be represented as a queueing diagram
  - Queueing diagram represents queues, resources, flows
  - We will discuss actual scheduling algorithms later

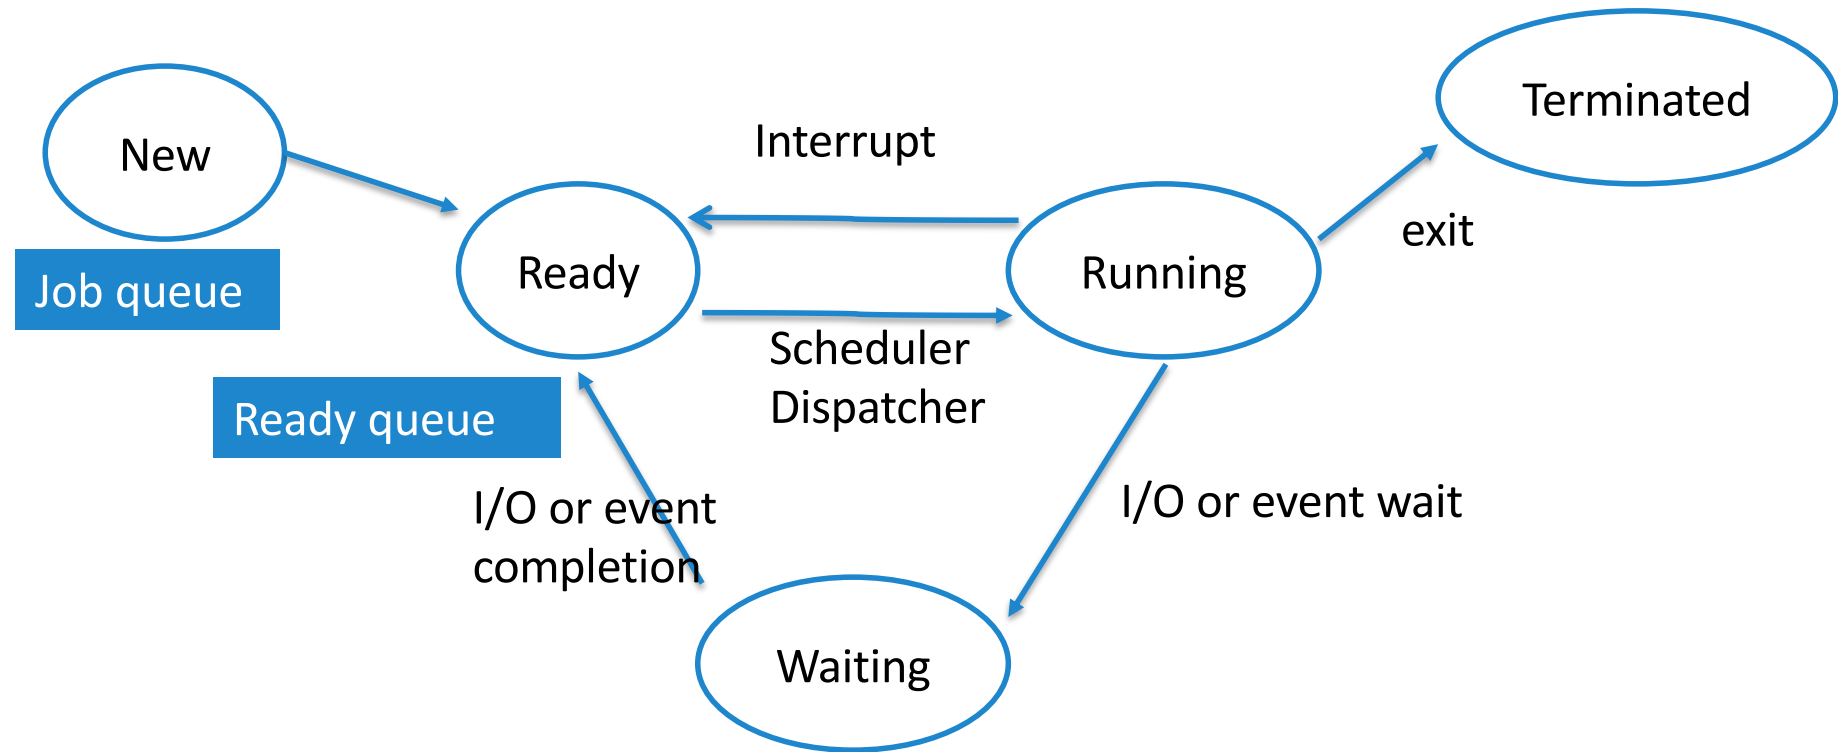
# Recap: Process state diagram



# Recap: Process state diagram



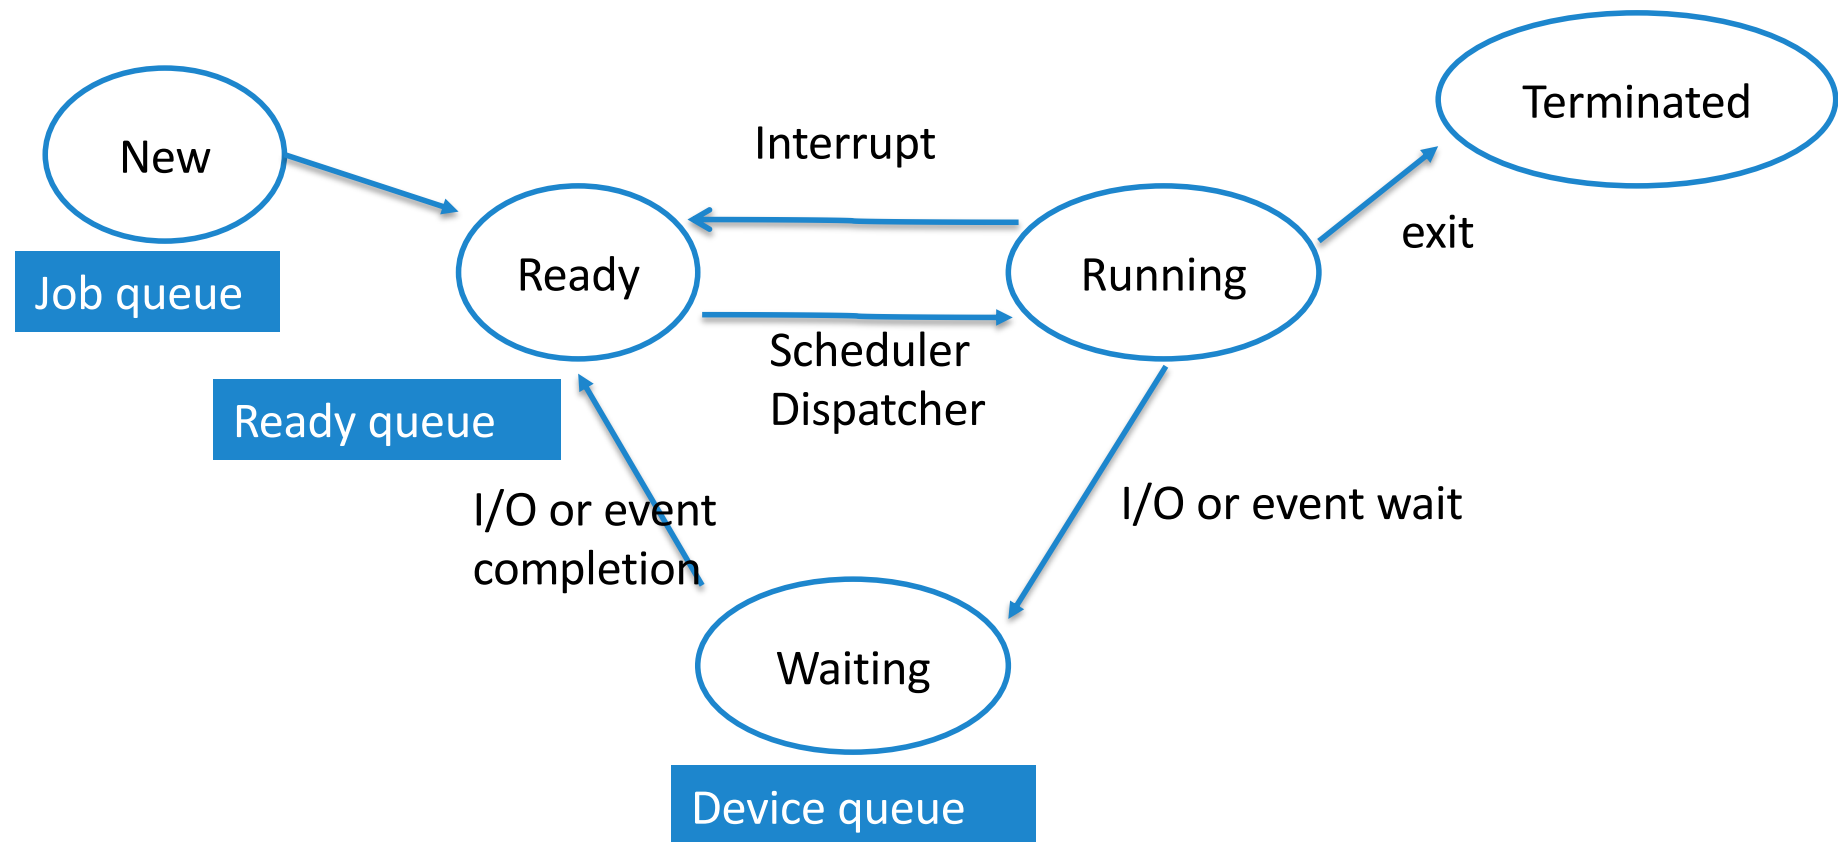
# Recap: Process state diagram



- The **process scheduler** selects a process from the ready queue for execution on the CPU
- **Dispatcher**: The kernel process that assigns the CPU to a process selected by the scheduler



# Recap: Process state diagram

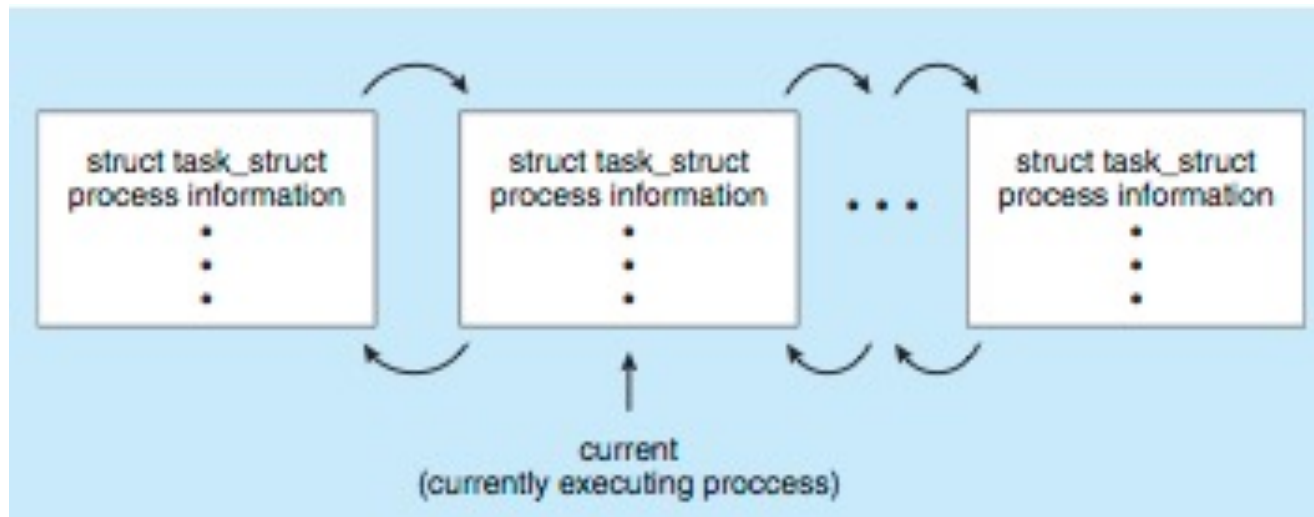


# Process queue representation in Linux

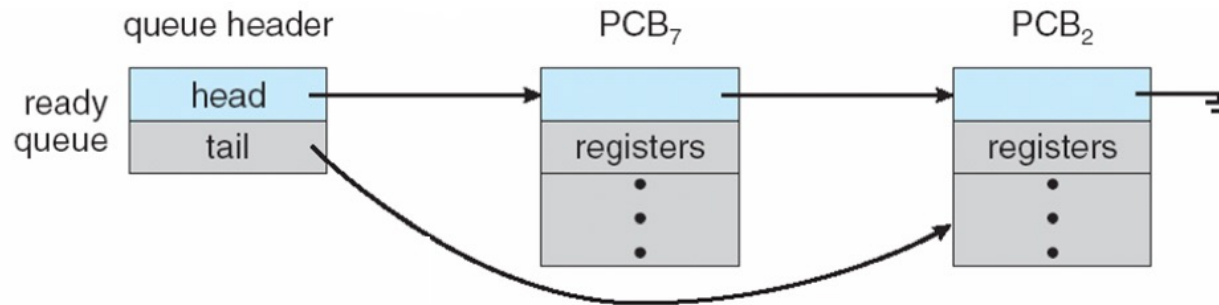
## Represented by the C structure `task_struct`

```
pid_t pid; // process identifier
long state; // state of the process
unsigned int time slice; // scheduling information
struct task_struct *parent; // this process's parent
struct task_struct *children; // this process's children
struct files_struct *files; // list of open files
struct mm_struct *mm; // address space of this pro
```

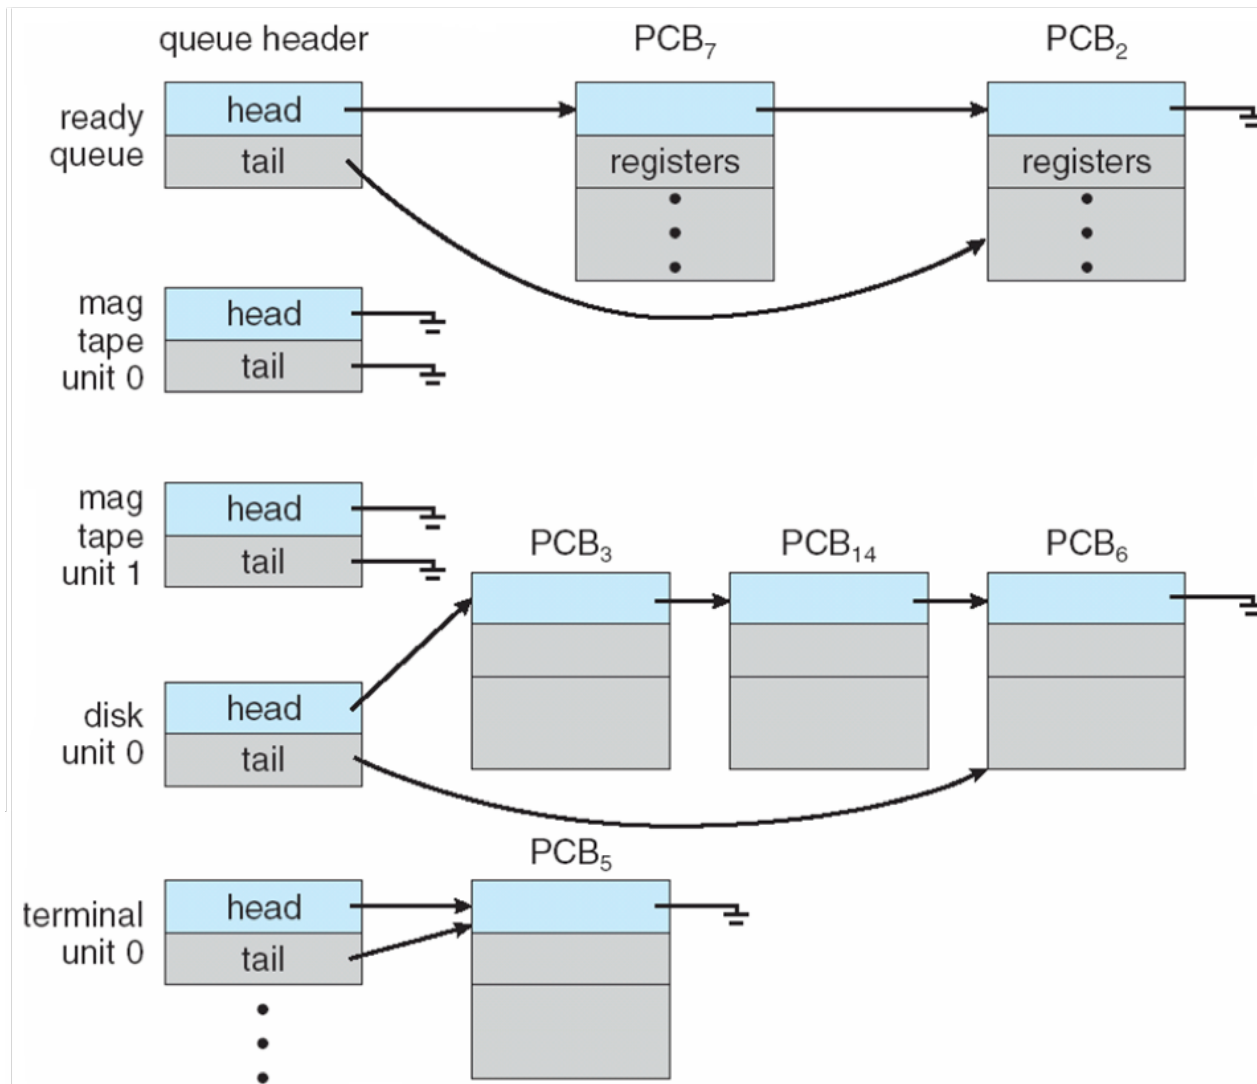
Doubly  
linked list



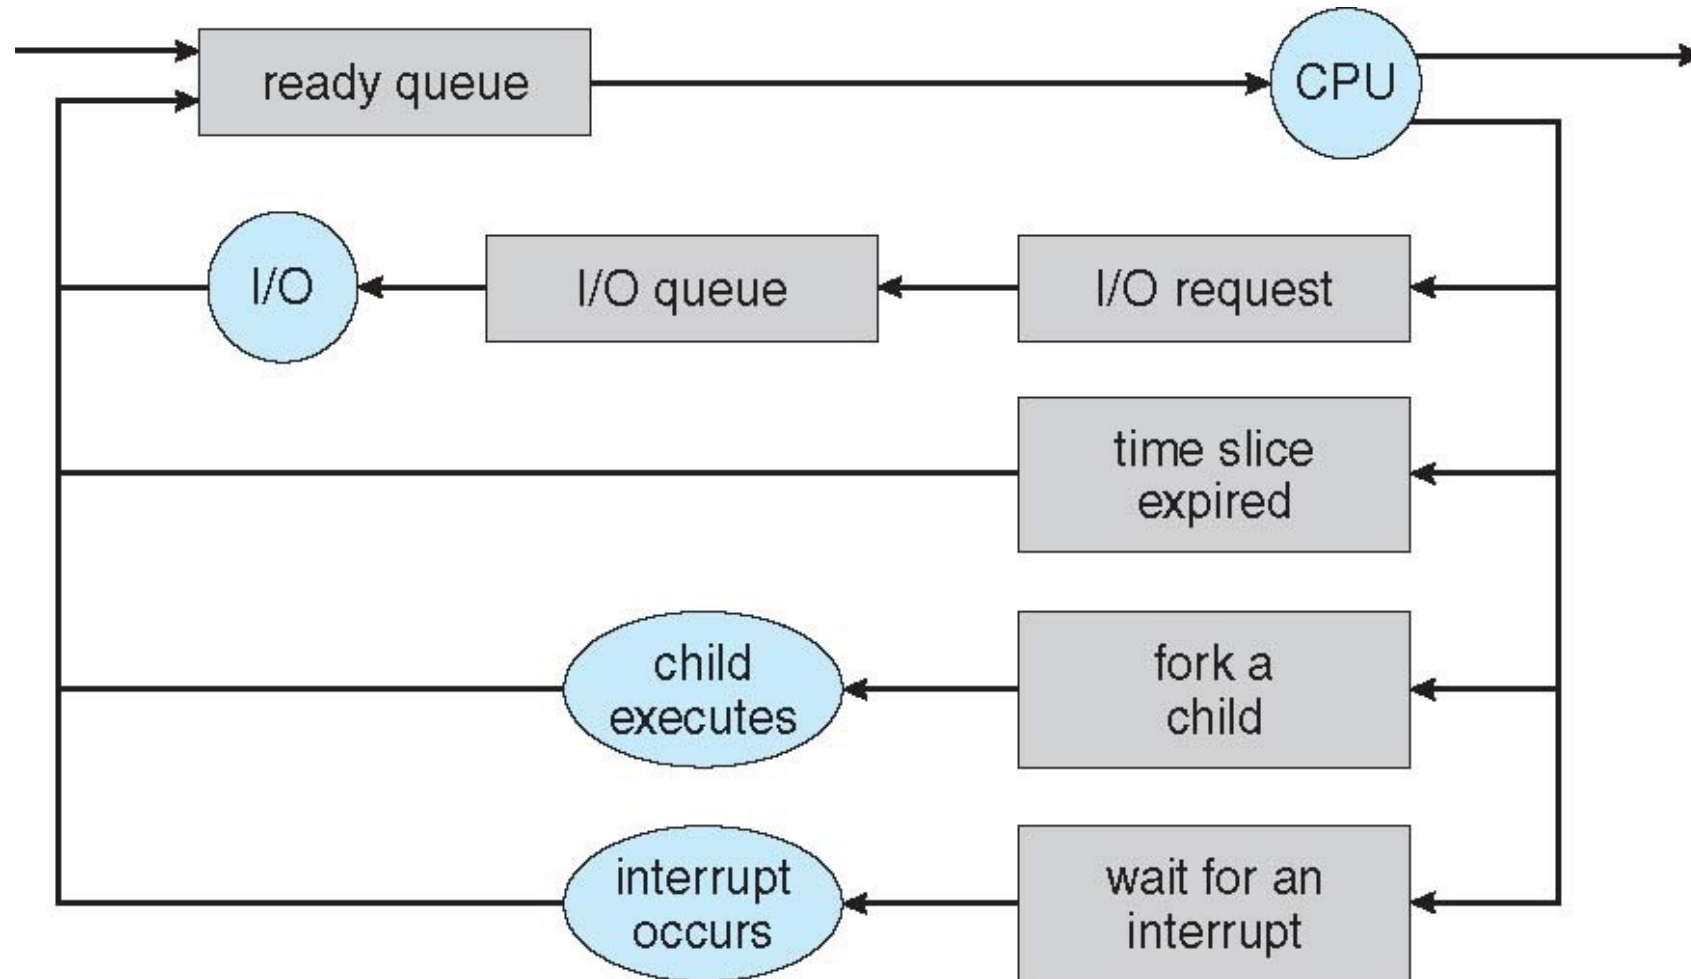
# Structure of process queues



# Structure of process queues



# Representation of process scheduling



# Operations on processes

# Process creation

- During execution a process may create several new processes
  - The process that creates a new process: **parent process**
  - The new process created: **child process**
  - Parent process create children processes, which, in turn, can create other processes, forming a tree of processes
  - Each process has a unique process identifier (pid)
  - Other than the first process (init), all other processes are created by **fork( ) system call**

# Process creation (contd.)

- Address space
  - Initially child is duplicate of parent
  - Child can later have a different program loaded into it
- In UNIX / Linux
  - `fork()` : creates a new process
  - `exec()`: replace new process's memory with new code

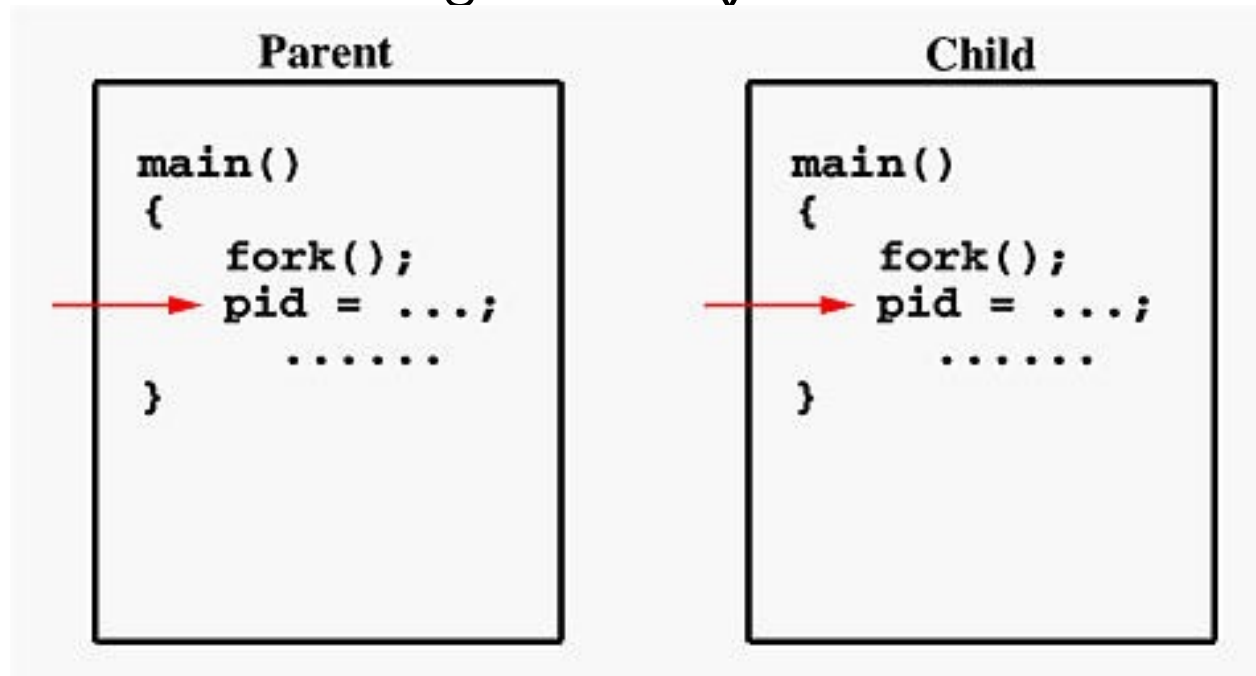


# fork( ) system call

- If `fork()` returns a negative value, the creation of a child process was unsuccessful.
- `fork()` returns a zero to the newly created child process.
- `fork()` returns a positive value, the *process ID* of the child process, to the parent.

# fork() system call

- If the call to `fork()` is executed successfully, the kernel will make an identical copy of the parent's address space, for use by the child.
- Both processes will start their execution at the next statement following the `fork()` call.



# Make a child do something different than the parent – method 1

## Parent

```
main()
{
    pid = 3456
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```

## Child

```
main()
{
    pid = 0
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```

# Make a child do something different than the parent – method 2

- `exec( )` system call and its variants
  - Replaces the calling process (that calls `exec`) with another program
  - Calling process terminated; the specified program is loaded in the same address space and then executed
  - New process image has same process ID (pid), same file descriptors, etc
  - Several variants – `execl`, `execle`, `execlp`, `execv`, `execve`, `execvp`

# Process waiting for a child

- A parent process can wait for a child to complete
  - Parent indicates to the kernel that it wants to wait, using the `wait()` system call
  - Parent execution is suspended; will be "ready" only after child terminates
  - If a child has already terminated, then the call returns immediately. Otherwise, block until a child terminates

# Process waiting for a child

- A parent process can wait for a child to complete
  - Parent indicates to the kernel that it wants to wait, using the `wait()` system call
  - Parent execution is suspended; will be "ready" only after child terminates
  - If a child has already terminated, then the call returns immediately. Otherwise, block until a child terminates
- Two variants
  - `wait()` - suspends execution of the current process until one of its children terminates
  - `waitpid(pid)` - suspends execution of the current process until the child specified by *pid* terminates

# Process termination

- A child process executes last statement
  - `exit()` call for deleting the process
  - Return status data to parent via `wait()` if any parent is waiting for the termination of this process
  - Deallocate the resources

# Process termination

- A child process executes last statement
  - `exit()` call for deleting the process
  - Return status data to parent via `wait()` if any parent is waiting for the termination of this process
  - Deallocate the resources

## Child process

.  
.

`exit(2) // Exit with status code`

## Parent process

`pid_t pid;`

`int status;`

.

`pid = wait (&status) // pid of  
terminated child`



# Process creation and exec: example

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) {
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) {
        execlp("/bin/ls", "ls", NULL);
    }
    else {
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

Annotations in the code:

- Error condition**: Points to the `if (pid < 0)` block.
- child process**: Points to the `else if (pid == 0)` block.
- parent process**: Points to the `else` block.

# Process termination: Special cases

- In some OS
  - All child must terminate when a process terminates
  - Cascading termination: All children, grandchildren etc. must be terminated when a process terminates
  - OS takes care of this cascade
- Combinations of `exit()` and `wait()`
  - If no parent is waiting, then **zombie process**
  - If parent terminated without invoking `wait()` then **orphan process**

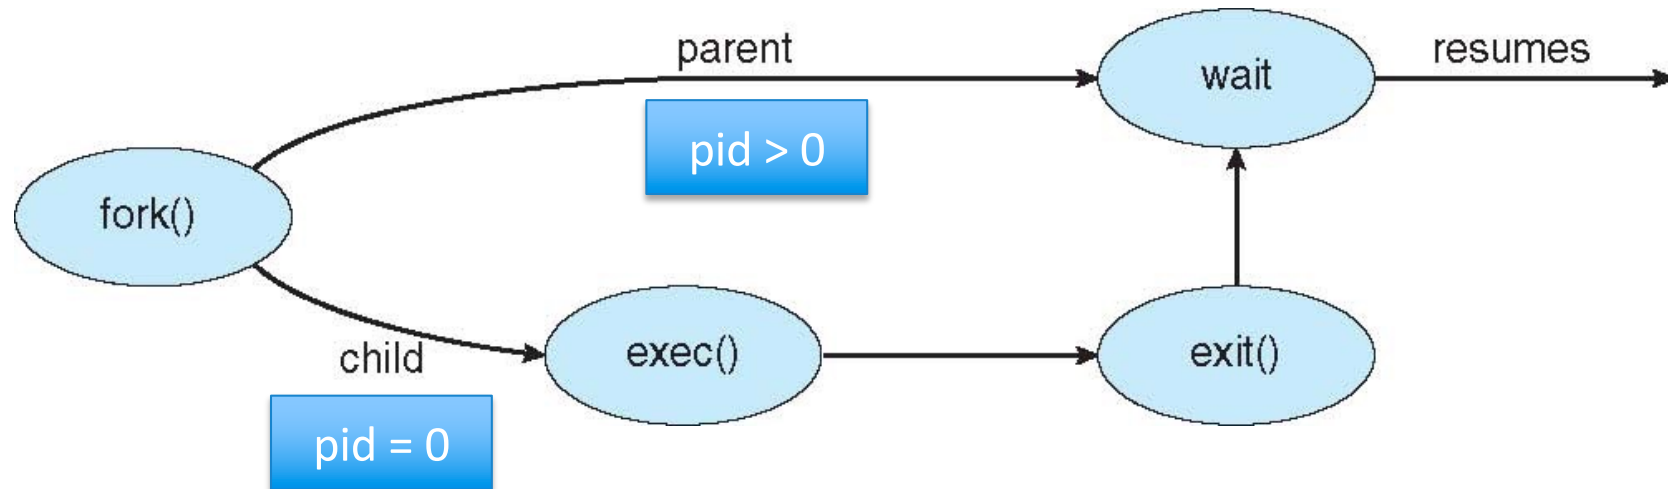
# Zombie and orphan process

- Zombie process
  - A process that has terminated, but whose parent has not (yet) called `wait()`
  - All processes move to this state when they terminate, and remain there until parent calls `wait()`
  - Entry in process table removed only after calling `wait()`

# Zombie and orphan process

- Zombie process
  - A process that has terminated, but whose parent has not (yet) called `wait()`
  - All processes move to this state when they terminate and remain there until parent calls `wait()`
  - Entry in process table removed only after calling `wait()`
- Orphan process
  - parent terminated without invoking `wait()`
  - Immediately “init” process assigned as parent
  - “init” periodically invokes `wait()`

# Usual workflow between parent & child



- Workflow can be different from what is shown
  - E.g., parent may or may not wait for the child
  - E.g., child may or may not use `exec()`

# Inter Process Communication

# Inter-process communication (IPC)

- Processes executing concurrently in OS may be independent or cooperating
- Cooperating process
  - Affect or be affected by other processes
  - Can share data
  - Speed-up in computation
  - Design can be modular

# Inter-process communication (IPC)

- Processes executing concurrently in OS may be independent or cooperating
- Cooperating process
  - Affect or be affected by other processes
  - Can share data
  - Speed-up in computation
  - Design can be modular
- Cooperating processes need IPC
  - Several types of IPC support provided by OS




# Inter-process communication (IPC)

- Ways to do IPC
  - 1: shared memory - `shmget()`, `shmcat()`, `shmaddr()`, `shmat()`, `shmdt()`, `shmctl()`
  - 2: message passing (pipe) - `pipe()`, `read()`, `write()`, `close()`
  - 3: message passing (named pipe) - `mkfifo()`, `read()`, `write()`, `close()`
  - 4: over network - RPC or Remote Procedure Call, sockets

**Shared memory system**

# Schematic for shared memory



Require the communicating processes to establish a region of shared memory

Relevant system calls in Linux / Unix

`shmget( )`

`shmat( )`

`shmdt( )`

`shmctl( )`

# Let's check the function calls

```
char *myseg;
```

```
key_t key; int shmid;
```

```
key = 235; // some unique id
```

```
shmid = shmget(key, 250, IPC_CREAT | 0666);
```

```
myseg = shmat(shmid, NULL, 0);
```

```
.
```

```
.
```

```
shmdt(myseg);
```

```
.
```

```
.
```

```
shmctl(shmid, IPC_RMID, NULL);
```

# Let's check the function calls

```
char *myseg;
```

```
key_t key; int shmid;
```

```
key = 235; // some unique id
```

Create shared memory segment

```
shmid = shmget(key, 250, IPC_CREAT | 0666);
```

```
myseg = shmat(shmid, NULL, 0);
```

```
•
```

```
•
```

```
shmdt(myseg);
```

```
•
```

```
•
```

```
shmctl(shmid, IPC_RMID, NULL);
```

# Let's check the function calls

```
char *myseg;
```

```
key_t key; int shmid;
```

```
key = 235; // some unique id
```

Create shared memory segment

```
shmid = shmget(key, 250, IPC_CREAT | 0666);
```

```
myseg = shmat(shmid, NULL, 0);
```

Attach the memory segment to  
the address space of this process

```
.
```

```
.
```

```
shmdt(myseg);
```

```
.
```

```
.
```

```
shmctl(shmid, IPC_RMID, NULL);
```

# Let's check the function calls

```
char *myseg;
```

```
key_t key; int shmid;
```

```
key = 235; // some unique id
```

Create shared memory segment

```
shmid = shmget(key, 250, IPC_CREAT | 0666);
```

```
myseg = shmat(shmid, NULL, 0);
```

Attach the memory segment to  
the address space of this process

```
.
```

```
.
```

```
shmdt(myseg);
```

Detach the memory segment from the  
address space of this process

```
.
```

```
.
```

```
shmctl(shmid, IPC_RMID, NULL);
```

# Let's check the function calls

```
char *myseg;
```

```
key_t key; int shmid;
```

```
key = 235; // some unique id
```

Create shared memory segment

```
shmid = shmget(key, 250, IPC_CREAT | 0666);
```

```
myseg = shmat(shmid, NULL, 0);
```

Attach the memory segment to  
the address space of this process

```
.
```

```
.
```

```
shmdt(myseg);
```

Detach the memory segment from the  
address space of this process

```
.
```

```
.
```

```
shmctl(shmid, IPC_RMID, NULL);
```

Mark the segment to be destroyed



# Example: Producer consumer problem

- A producer process produces information that is consumed by the consumer process
  - Compiler produces assembly code consumed by assembler
  - Program produces lines to print, print spool consumes
  - The information is written to / read from a buffer

# Example: Producer consumer problem

- A producer process produces information that is consumed by the consumer process
  - Compiler produces assembly code consumed by assembler
  - Program produces lines to print, print spool consumes
  - The information is written to / read from a buffer
- Two variants
  - Bounded buffer: producer waits when buffer is full, consumer waits when buffer is empty
  - Unbounded buffer: consumer waits when buffer is empty

# Producer consumer solution with bounded buffer

- Shared data: implemented as a circular array

```
#define BUFFER_SIZE 10
```


```
typedef struct {  
    ... // info to be shared  
} item;
```

```
item buffer[BUFFER_SIZE];
```

```
int in = 0;
```

```
int out = 0;
```

These variables reside  
in shared memory




# Key ideas

- Circular buffer
  - Index in: the next free position to write to
  - Index out: the next filled position to read from
- To check buffer full or empty:
  - Buffer empty:  $in == out$
  - Buffer full:  $in + 1 \% BUFFER\_SIZE == out$ 
    - This scheme allows at most  $BUFFER\_SIZE - 1$  items in the buffer

# Pseudo code

```
while (true) {  
    /* Produce an item */  
    while (((in + 1) % BUFFER_SIZE) == out)  
        ; /* do nothing -- no free buffers */  
    buffer[in] = newProducedItem;  
    in = (in + 1) % BUFFER_SIZE;  
}
```

Producer



```
while (true) {  
    while (in == out)  
        ; // do nothing -- nothing to consume  
  
    // remove an item from the buffer  
    itemToConsume = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    return itemToConsume;  
}
```

Consumer

Solution is correct, but can only use  
BUFFER\_SIZE-1 elements

# Better utilization of buffer space

- Circular buffer (shared memory)
- Suppose that we want to use all buffer space:
  - an integer count: the number of filled buffers (shared memory)
  - Initially, count is set to 0.
  - incremented by producer after it produces a new buffer
  - decremented by consumer after it consumes a buffer.

# Better utilization of buffer space: Pseudo code

## Producer

```
while (true) {  
    /* produce an item  
       and put in nextProduced */  
    while (count == BUFFER_SIZE)  
        ; // do nothing  
  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

## Consumer

```
while (true) {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
  
    /* consume the item in  
       nextConsumed */  
}
```

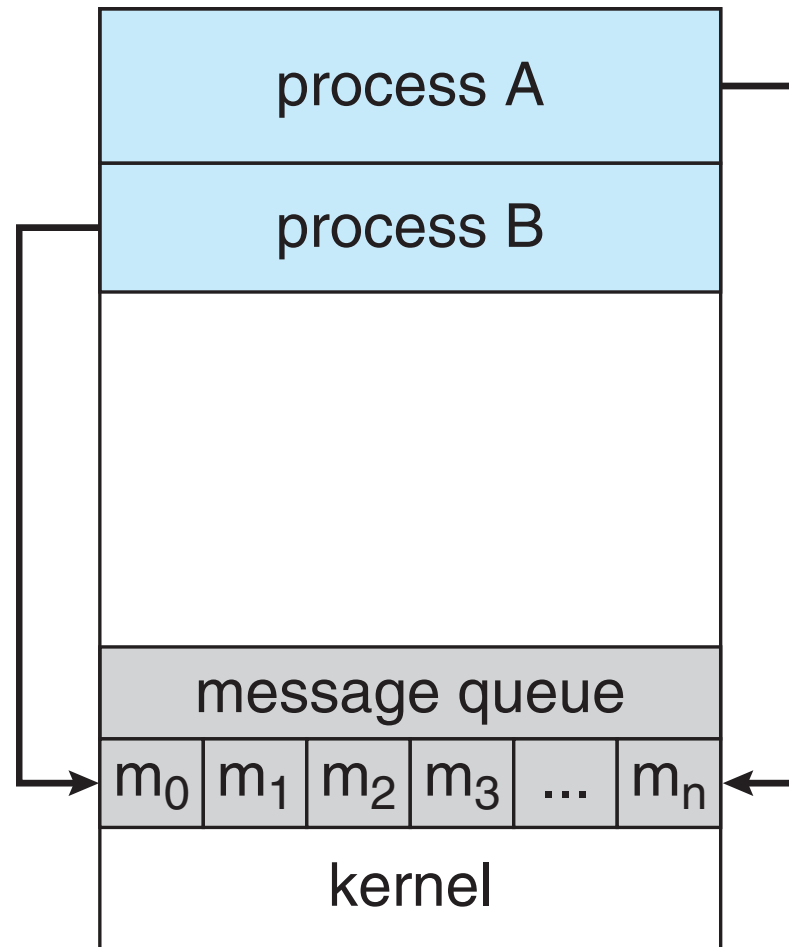
**Message passing system**



# Basics of message passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - **send**(*message*)
  - **receive**(*message*)
- The *message* size is either fixed or variable

# Communication model



# Ways for message passing

- Pipes
- Named pipes

# Pipes

- Acts as a medium to allow two processes to communicate
  - Communication can be uni/bi-directional
  - Must there exist a relationship (i.e., parent-child) between the communicating processes?
  - Can the pipes be used over a network?

# Ordinary pipes

- A message passing medium between related processes
  - Cannot be accessed from outside the process
  - Typically, a parent process creates a pipe and uses it to communicate with its child process
  - Pipes behave like FIFO queues
  - Read-write in pipe == producer-consumer
  - Producer writes to one end (the write-end of pipe)
  - Consumer reads from the other end (the read-end of pipe)
  - Unidirectional

# Producer consumer in pipes

## Producer

```
message next_produced;  
  
while(TRUE){  
    <produce and put data in next produced>  
    ...  
    send( next_produced);  
}
```

## Consumer

```
message next_consumed;  
  
while(TRUE){  
    receive(next_consumed);  
    ...  
    <data in next consumed>  
}
```

# Named pipes

- Accessed as files by processes
  - No parent-child relation is necessary
  - Still behave like FIFO queues (even called fifo)
  - Several processes can use the named pipe

# Named pipes

- Accessed as files by processes
  - No parent-child relation is necessary
  - Still behave like FIFO queues (even called fifo)
  - Several processes can use the named pipe

```
char* myfifo = '/tmp/myfifo';  
mkfifo (myfifo, 0666); // creates the fifo or named pipe (a special file)  
...  
fd = open(myfifo, O_WRONLY); // Process A  
write(fd, ...); // Process A  
close(fd); // Process A  
...  
fd = open(myfifo, O_RDONLY); // Process B  
read(fd, ...); // Process B  
close(fd); // Process B
```



# Named pipes

- Once you have created a FIFO special file / named pipe, any process can open it for reading or writing, in the same way as an ordinary file.
- However, it has to be open at both ends simultaneously before you can proceed to do any input or output operations on it.

# Finally, for communication of two processes over network

- Sockets API
- Remote procedure call

# Summary

- What is a process?
  - Structure of a process
  - Process states
  - Process control block
  - Context switch
- Queues for process scheduling
  - Ready queues, event queues, queueing diagram
- How does two processes talk?
  - Shared memory, pipe, named pipe